

# Submodular Maximisation

Andreas Göbel

Hasso Plattner Institute


2020

# Examples of AI tasks (I)

- **Influence maximisation on a social network**

Select a subset of the most influential users on the network.

e.g. viral marketing, personalised recommendation, selecting influential tweeters

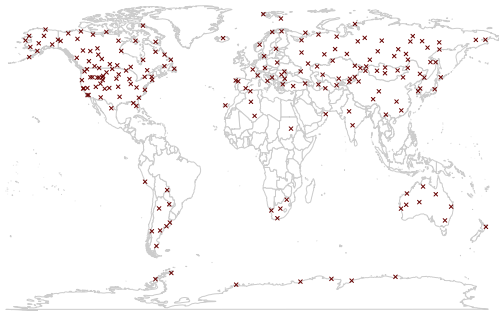


## Examples of AI tasks (II)

- Experimental design

Design an experiment in a way that optimises the accuracy of information provided by this experiment.

e.g. statistical tests, sensor placement



# Examples of AI tasks (III)

- Automatic summarisation

Shorten a set of data that represents the most important or relevant information within the original content.

e.g. text, images and video



# Submodular set functions

A set function  $f$  defined on a ground set  $\Omega$  takes as input a subset  $S \subset \Omega$  and return a real value  $f(S) \in \mathbb{R}$ .

## Definition (Submodular set function)

A function  $f : \mathcal{P}(\Omega) \rightarrow \mathbb{R}$  is submodular if for all sets  $A, B \subseteq \Omega$

$$f(A) + f(B) \geq f(A \cap B) + f(A \cup B)$$


Submodularity captures the notion of diminishing returns.

For sets  $A \subseteq B \subseteq \Omega$  and for each  $x \notin B$ , we have

$$f(A \cup \{x\}) - f(A) \geq f(B \cup \{x\}) - f(B)$$

# Video summarisation

We want to summarise a video consisting of  $n$  frames.

- For each frame  $i$  we can extract a feature vector  $\vec{x}_i$  using a neural network.
  - ▶ Examples: colour, luminosity, number of faces, SIFT.
- We generate an  $n \times n$  matrix  $M$ , where  $M_{i,j}$  expresses how similar the feature vectors  $\vec{x}_i$  and  $\vec{x}_j$  are.
- To select the most diverse frames of the video we have to find the subset of frames  $S \subseteq \{1, \dots, n\}$  maximizing the value

$$f(S) = \text{Det}(M|_S),$$

where  $M|_S$  is the submatrix of  $M$  restricted on the rows and columns of  $S$ .


- It turns out that the function  $f$  is submodular.
- Furthermore the function is monotone, i.e. for  $A \subseteq B$ ,  $f(A) \leq f(B)$ .

# Side constraints

In such applications the computational task can be abstracted as follows: Given a submodular function  $f : \mathcal{P}(\Omega) \rightarrow \mathbb{R}$  find the set  $S \subseteq \Omega$  that maximises  $f(S)$ .

Depending on the application maximising the function  $f$  might require further restrictions on the set  $S$ .

- **Cardinality constraint.**  $|S| \leq k$ .
- **Partition constraint.**  $\Omega$  can be partitioned in  $\{\Omega_i\}_{i \leq m}$ , we require  $|S \cap \Omega_i| \leq k_i$ . (**Matroid constraint**)
- **Knapsack constraint.** Each element  $x \in \Omega$  has a weight  $w(x)$ , we require  $\sum_{x \in S} w(x) \leq W$ .



# Computational aspects

- Computing the value of the function  $f$  often requires a lot of computational resources (time).
- Grey-box complexity: the efficiency of an algorithm is measured as a number of (oracle) evaluations of the function  $f$ .

Maximising a submodular function is **NP-complete** (MaxCut is a special case). It is unlikely to find an exact solution with at most polynomial number of function evaluations in terms of  $n = |\Omega|$ .

Efficient approximation algorithms work well in theory and in practice.



# A simple algorithm

The Greedy algorithm builds a solution by iteratively adding the element that yields the largest marginal gain.

- Start with the empty set  $S_0 = \emptyset$ .
- At each time step  $t$  find the element  $x$  maximizing  $f(S_t \cup \{x\}) - f(S_t) > 0$  and add  $x$  to the current solution.
- Stop when no element can be added.  
Due to constraints on  $S$  or no element gives positive marginal gain.

$O(n^2)$  evaluations of  $f$ .

Greedy is often used in practice due to its simplicity. In theory to achieve better theoretical approximation guarantees more elaborate algorithms are used.


# Approximation guarantees

A  $1/2$ -approximation linear-time algorithm is known for the unconstrained problem. Constraints make the problem more complicated.

Commonly in applications the function  $f$  is monotone.

- The greedy gives a  $(1 - 1/e)$ -approximation guarantee under a cardinality constraint.
- $(1 - 1/e)$ -approximation algorithms are known for a matroid constraint.  $(1/2)$ -approximation for the greedy.
- $(1 - 1/e)$ -approximation algorithms are known for a knapsack constraint. No approximation guarantees for greedy.

Greedy is usually the practitioners choice as it performs extremely well in real world instances.



# Further computational variants

## Streaming

- Can you summarise data “on the fly”?
- Streaming algorithms require only a single pass through the data.
- The greedy algorithm requires multiple passes.
- For the cardinality constraint a greedy streaming algorithm exchanging the element in the solution achieves a  $1/2$ -approximation guarantee.

## Adaptive complexity

- Function evaluations can often be done in parallel.
- For greedy we require  $O(n)$  adaptive rounds.
- With a more clever algorithm we can reduce this to  $O(\log^3 n)$  while maintaining the same approximation guarantees.